

Attention Is All You Need

We propose the Transformer, a model architecture based solely on attention mechanisms, dispensing with recurrence and convolutions entirely.

Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring less time to train.

The Transformer uses stacked self-attention and point-wise, fully connected layers for both the encoder and decoder.